

Utkarsh Mishra

utkarshmishra2k2@gmail.com
+91 94284 68492

linkedin.com/in/utkarsh-mishra-1aa340227
github.com/Utkarshmishra2k2

Results-driven AI professional with hands-on experience building production-grade systems at government and tech organizations. Specialized in LLM architecture, RAG pipelines, statistical modeling, and full-stack development. Proven ability to translate complex ML concepts into scalable, user-facing products with measurable impact on performance, accuracy, and business outcomes.

Skills

- **Language:** Python, R, SQL, Linux OS
 - **Machine Learning:** Supervised Learning (Regression, Classification, SVM, k-NN), Unsupervised Learning (Clustering, Factor Analysis), Tree-based Methods (Decision Trees, Random Forests), Ensemble Methods, Feature Engineering, Regularization (L1/L2), Hyperparameter Tuning
 - **Generative AI & LLMs:** Prompt Engineering, RAG, Finetuning, LangChain, LangGraph, Transformer Models, Ollama, Vector Databases, Embedding Models, Hallucination Mitigation
 - **Statistics:** Hypothesis Testing, Sampling, Distribution Theory, Regression, Classification, Factor Analysis
 - **Data & ML Tools:** Pandas, NumPy, Scikit-learn, TensorFlow, Seaborn, Power BI, Matplotlib
 - **Backend & Full-Stack:** FastAPI, Python web services, JWT authentication, REST APIs
-

Professional Experience

International Financial Services Centres Authority (IFSCA), Govt of India:

- AI – Large Language Models (LLM) | Jul 2025 – Present | Gandhinagar
- Built an agentic AI-driven enterprise service desk using FastAPI, where LLM agents autonomously classify user intent, route workflows, and manage multi-step conversations for meetings, equipment, and IT tickets. Implemented strict trust boundaries, confidence-gated orchestration, RBAC, and audited, production-grade integrations.
 - Built IFSCA IntelliChat, a production RAG platform for 100+ concurrent users querying 500+ regulatory PDFs with sub-second latency and precise citations.
 - Architected full-stack system: Python, LangChain/LangGraph, Llama 3.1 (Ollama), FAISS, FastAPI, React (Vite/Tailwind), JWT auth, audit logs, offline ops.
 - Cut hallucinations 35% via dynamic top-K retrieval and regulatory embedding fine-tuning.
 - Applied LASSO logistic regression & Bayesian models on AI adoption surveys for GIFT City policy insights.

Dynamite Technology: Data Analyst | Mar 2024 – Apr 2024 | Navi Mumbai

- Built recommendation system using TF-IDF + cosine similarity, achieving 0.78 precision@10 on holdout test set.
- Optimized Pandas data pipeline, cutting preprocessing time 45% via vectorized ops and feature selection for production.

Projects

Focus on Vision: A Statistical Study of Eye:

Tech Stack: Python, TensorFlow, Power BI, Streamlit

Impact: 1,400+ survey participants, 86.9% accuracy

- o Designed and deployed a 60-question survey capturing demographics, screen time, and genetic factors; built Power BI dashboards visualizing vision-risk correlations.
- o Engineered features through EDA, handled missing value imputation, and applied SMOTE to address class imbalance before model training.
- o Built a dual-model architecture combining an interpretable logistic regression model (Recall: 0.84) with a high-performance TensorFlow neural network (Accuracy: 86.9%, ROC-AUC: 0.81).
- o Deployed Streamlit app for real-time personalized risk assessment with confidence scores and feature importance.

AI-Powered Data Analytics Platform:

Tech Stack: Gemini LLM, Python, Streamlit, DataStax Vector DB

Impact: Dynamic code generation

- o Architected platform enabling natural-language queries to automatically generate executable Python code for data analysis and visualization without manual coding.
- o Fine-tuned domain-specific transformer models to improve contextual precision and reduce hallucination on financial and healthcare datasets.
- o Built stateful Streamlit UI with DataStax backend for session management, embedding storage, and conversational analytics—supporting multi-turn dialogue and chat history.

Education

- **M.Sc in Statistics** | CGPA: 8.90 / 10 | University of Mumbai | 2023 - 2025
 - **B.Sc in Statistics** | CGPA: 9.40 / 10 | University of Mumbai | 2020 - 2023
-

Certifications

- **Generative AI with Large Language Models** - Coursera
- **Data Analytics using R Software & Excel** - Ramniranjan Jhunjhunwala College
[Value Added Courses]
- **TCS iON National Qualifier Test** - IT Programming